

Mining Data Streams (Part 1)

Ahmad Al-Shishtawy and Vladimir Vlassov

Acknowledgement: Some of the slides are adopted from the slide provided at <http://www.mmids.org> for the book “Mining of Massive Datasets” by Jure Leskovec, Anand Rajaraman, Jeff Ullman, Stanford University



Outline

- Part 1 (this lecture)
 - The stream data model
 - Sampling data in a stream
 - Filtering streams
- Part 2 (next lecture: counting algorithms on streams)
 - Counting distinct elements
 - Computing moments
 - Queries over a (long) sliding window: Counting ones
 - Decaying Windows: Counting itemsets



Data Streams – Infinite Data

In many data mining situations, we do not know the entire data set in advance

Stream Management is important when the input rate is controlled **externally**:

- Google search queries
- Twitter or Facebook status updates

We can think of the **data** as **infinite** and **non-stationary** (the distribution changes over time)



The Stream Model

Input **elements** enter at a rapid rate, at one or more input ports (i.e., **streams**)

- **We call elements of the stream tuples**

The system cannot store the entire stream accessibly

Q: How do you make critical calculations about the stream using a limited amount of (secondary) memory?

SGD – Example of a Streaming Algorithm

Stochastic Gradient Descent (SGD) is an example of a stream algorithm


In Machine Learning we call this: Online Learning

- Allows for modeling problems where we have a continuous stream of data
- We want an algorithm to learn from it and slowly adapt to the changes in data

Idea: Do slow updates to the model

- **SGD** (SVM, Perceptron) makes small updates
- **So:** First train the classifier on training data.
- **Then:** For every example from the stream, we slightly update the model (using small learning rate)

General Stream Processing Model



Two forms of query

1. **Ad-hoc queries**: Normal queries asked one time about streams.
 - **Example**: What is the maximum value seen so far in stream S?
2. **Standing queries**: Queries that are, in principle, asked about the stream at all times.
 - **Example**: Report each new maximum value ever seen in stream S.



Problems on Data Streams (1)

Types of queries one wants on answer on a data stream:

(to be considered in this lecture)

- **Sampling data from a stream**
 - Construct a random sample
- **Filtering a data stream**
 - Select elements with property x from the stream

Problems on Data Streams (2)

Types of queries one wants on answer on a data stream: (to be considered in next lecture)

- **Queries over sliding windows**
 - Number of items of type x in the last k elements of the stream
- **Counting distinct elements**
 - Number of distinct elements in the last k elements of the stream
- **Estimating moments**
 - Estimate avg./std. dev. of last k elements
- **Finding frequent elements**



Applications (1)

Mining query streams

- Google wants to know what queries are more frequent today than yesterday

Mining click streams

- Yahoo! wants to know which of its pages are getting an unusual number of hits in the past hour
 - Often caused by annoyed users clicking on a broken page.

Mining social network news feeds

- E.g., look for trending topics on Twitter, Facebook



Applications (2)

Sensor Networks

- Many sensors feeding into a central controller

Telephone call records

- Data feeds into customer bills as well as settlements between telephone companies

IP packets monitored at a switch

- Gather information for optimal routing
- Detect denial-of-service attacks

Sampling from a Data Stream: Sampling a fixed proportion

As the stream grows the sample also gets bigger

Sampling from a Data Stream

Since **we can not store the entire stream**, one obvious approach is to store a **sample**

Two different problems:

- (1) Sample a **fixed proportion** of elements in the stream (say 1 in 10)
- (2) Maintain a **random sample of fixed size** over a potentially infinite stream
 - At any “time” k we would like a random sample of s elements
 - **What is the property of the sample we want to maintain?**
For all time steps k , each of k elements seen so far has equal probability of being sampled

Sampling a Fixed Proportion

Problem 1: Sampling fixed proportion

Scenario: Search engine query stream

- **Stream of tuples:** (user, query, time)
- **Answer questions such as:** How often did a user run the same query in a single days
- Have space to store $1/10^{\text{th}}$ of query stream

Naïve solution:

- Generate a random integer in $[0..9]$ for each query
- Store the query if the integer is **0**, otherwise discard

Problem with Naïve Approach

Simple question: What fraction of queries by an average search engine user are duplicates?

- Suppose each user issues x queries once and d queries twice, i.e., in total of $(x + 2d)$ queries
 - **Correct answer is** $d/(x + d)$
- **Proposed solution: Let's keep 10% of the queries**
 - Sample will contain $x/10$ of the singleton queries and $2d/10$ of the duplicate queries at least once
 - But only $d/100$ pairs of duplicates: $d/100 = 1/10 \cdot 1/10 \cdot d = 0,01d$
 - Of d “duplicates” $18d/100$ appear exactly once: $((1/10 \cdot 9/10) + (9/10 \cdot 1/10)) \cdot d = 0,18d$
- **So the sample-based answer is**
$$\frac{0,01d}{0,1x + 0,01d + 0,18d} = \frac{d}{10x + 19d}$$



Solution: Sample Users

- Pick $1/10^{\text{th}}$ of **users** and take all their searches in the sample
- Use a hash function that hashes the user name or user id uniformly into 10 buckets

Generalized Solution

Stream of tuples with keys:

- Key is some subset of each tuple's components
 - e.g., tuple is (user, search, time); key is **user**
- Choice of key depends on application

To get a sample of a/b fraction of the stream:

- Hash each tuple's key uniformly into b buckets
- Pick the tuple if its hash value is at most a



Hash table with b buckets, pick the tuple if its hash value is at most a .

How to generate a 30% sample?

Hash into $b=10$ buckets, take the tuple if it hashes to one of the first $a=3$ buckets

Sampling from a Data Stream: Sampling a fixed-size sample

As the stream grows, the sample is of a fixed size

Maintaining a Fixed-size Sample

Problem 2: Fixed-size sample

Suppose we need to maintain a random sample S of size exactly s tuples

- E.g., main memory size constraint

Why? Don't know length of stream in advance

Suppose at time n we have seen n items

- Each item is in the sample S with equal prob. s/n

Maintaining a Fixed-size Sample (cont)

How to think about the problem: say the sample size $s = 2$

Stream: a x c y z k c d e g...

At $n = 5$, each of the first 5 tuples is included in the sample S with equal probability.

At $n = 7$, each of the first 7 tuples is included in the sample S with equal probability.

Impractical solution would be to store all the n tuples seen so far and out of them pick s at random

Solution: Fixed Size Sample

Algorithm (a.k.a. Reservoir Sampling)

- Store all the first s elements of the stream to S
- Suppose we have seen $n-1$ elements, and now the n^{th} element arrives ($n > s$)
 - With probability s/n , keep the n^{th} element, else discard it
 - If we picked the n^{th} element, then it replaces one of the s elements in the sample S , picked uniformly at random

Claim: This algorithm maintains a sample S with the desired property:

- After n elements, the sample contains each element seen so far with probability s/n

Proof: By Induction

We prove this by induction:

- Assume that after n elements, the sample contains each element seen so far with probability s/n
- We need to show that after seeing element $n+1$ the sample maintains the property
 - Sample contains each element seen so far with probability $s/(n+1)$

Base case:

- After we see $n=s$ elements the sample S has the desired property
 - Each out of $n=s$ elements is in the sample with probability $s/s = 1$

Proof: By Induction (cont)

Inductive hypothesis: After n elements, the sample S contains each element seen so far with probability s/n

Now element $n+1$ arrives

Inductive step: For elements already in S , probability that the algorithm keeps it in S :

$$\underbrace{\left(1 - \frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ discarded}} + \underbrace{\left(\frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ not discarded}} \underbrace{\left(\frac{s-1}{s}\right)}_{\text{Element in the sample not picked}} = \frac{n}{n+1}$$

- So, at time n , tuples in S were there with probability s/n
- At time $n \rightarrow n+1$, tuple stayed in S with prob. $n/(n+1)$
- So prob. tuple is in S at time $n+1$: $(s/n) \cdot (n/(n+1)) = \frac{s}{n+1}$

Filtering Data Streams



Filtering Data Streams

Another common process on streams is *selection*, or *filtering*.

- Each element of data stream is a tuple
- Tuples in the stream that meet a criterion are accepted.
 - Accepted tuples are passed to another process as a stream
 - Non-accepted tuples are dropped.



Filtering Data Streams

Each element of data stream is a tuple

Given a list of keys S

Determine which tuples of stream are in S

Obvious solution: Store keys in a hash table

- But suppose we **do not have enough memory** to store all of S in a hash table
 - E.g., we might be processing millions of filters on the same stream



Applications

Example: Email spam filtering

- We know 1 billion “good” email addresses
- If an email comes from one of these, it is **NOT** spam

Publish-subscribe systems

- You are collecting lots of messages (news articles)
- People express interest in certain sets of keywords
- Determine whether each message matches user’s interest

First Cut Solution (1)

Given a set of keys S that we want to filter

Create a **bit array B** of n bits, initially all **0s**

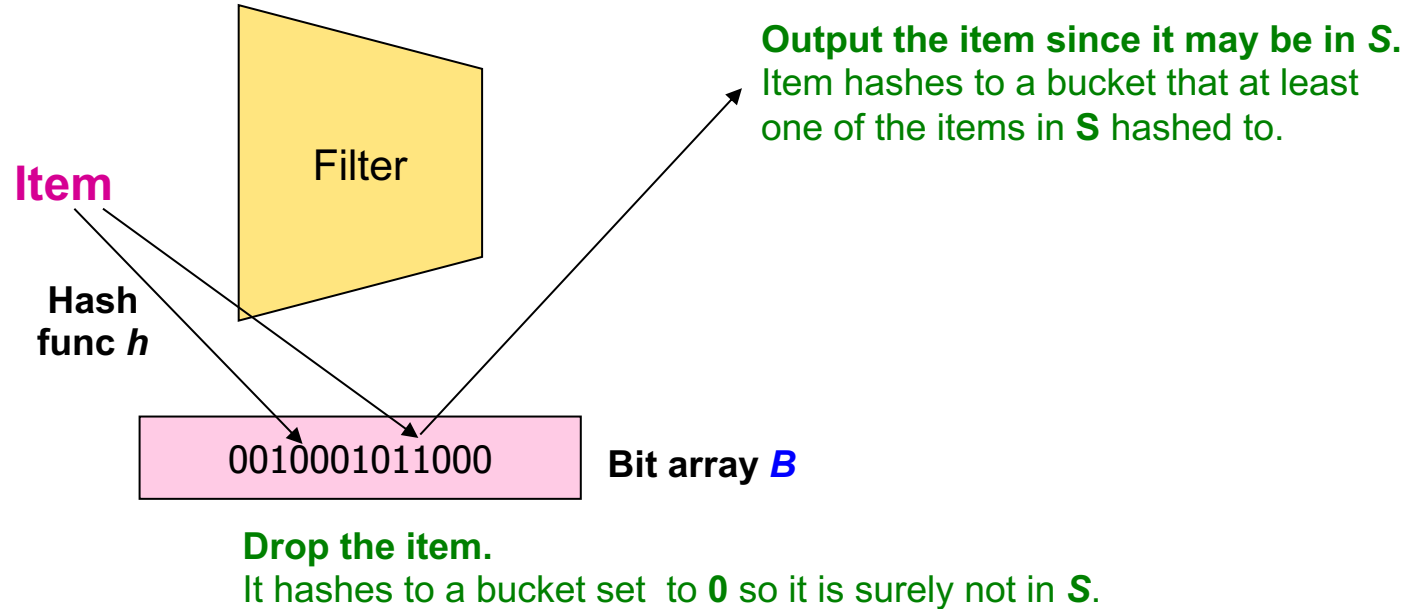
Choose a **hash function h** with range **$[0, n)$**

Hash each member of $s \in S$ to one of n buckets, and set a corresponding bit to **1**, i.e., **$B[h(s)] = 1$**

Hash each element a of the stream and output only those that hash to a bucket with its corresponding bit in B that was set to **1**

- **Output a if $B[h(a)] == 1$**

First Cut Solution (2)



Creates false positives but no false negatives

- If the item is in S we surely output it, if not we may still output it

First Cut Solution (3)

- $|S| = 1$ billion email addresses
 $|B| = 1\text{GB} = 8$ billion bits

If the email address is in S , then it surely hashes to a bucket that has the bit set to 1, so it always gets through (*no false negatives*)

Approximately $1/8$ of the bits are set to 1, so about $1/8^{\text{th}}$ of the addresses not in S get through to the output (*false positives*)

- Actually, less than $1/8^{\text{th}}$, because more than one address might hash to the same bit

Bloom Filter

A Bloom filter consists of:

1. An bit-array B of n bits, initially all 0's.
2. A collection of hash functions h_1, h_2, \dots, h_k .
Each hash function maps “key” values to n buckets, corresponding to the n bits of the bit-array.
3. A set S of m key values

The Bloom filter allows through all stream elements whose keys are in S

- It rejects stream elements whose keys are not in S

Bloom Filter (cont)

Initialization:

- Set \mathbf{B} to all 0s
- Hash each element $\mathbf{s} \in \mathbf{S}$ using each hash function h_i , set $\mathbf{B}[h_i(\mathbf{s})] = 1$ (for each $i = 1, \dots, k$)

Run-time:

- When a stream element with key \mathbf{x} arrives
 - If $\mathbf{B}[h_i(\mathbf{x})] = 1$ for all $i = 1, \dots, k$ then declare that \mathbf{x} is in \mathbf{S}
 - That is, \mathbf{x} hashes to a bucket set to 1 for every hash function $h_i(\mathbf{x})$
 - Otherwise discard the element \mathbf{x}

Analysis of Bloom Filtering: Throwing Darts

Analysis for the number of **false positives**

(assuming one hash function $k = 1$)

Consider: If we throw m darts into n equally likely targets, **what is the probability that a target gets at least one dart?**

In our case:

- **Targets** = bits/buckets
- **Darts** = hash values of items

Bloom Filter – Analysis

Bloom filter: bit-array **B** of n bits; k hash functions; m keys

What fraction of the bit vector **B are 1s?**

- Throwing $k \cdot m$ darts at n targets
- So fraction of 1s is $(1 - e^{-km/n})$

But we have k independent hash functions and we only let the element x to pass through **if all** k hash functions hash element x to a bucket of value 1

So, false **positive probability** = $(1 - e^{-km/n})^k$

Bloom Filter – Analysis (cont)

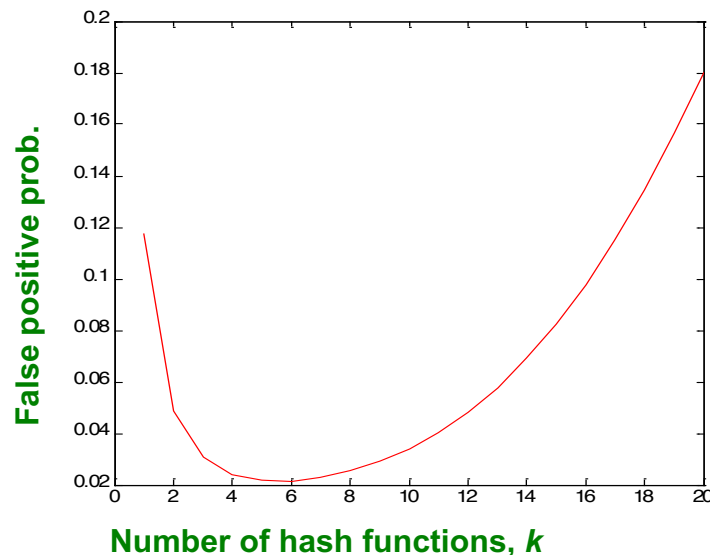
$m = 1$ billion, $n = 8$ billion

- $k = 1$: $(1 - e^{-1/8}) = 0.1175$
- $k = 2$: $(1 - e^{-1/4})^2 = 0.0493$

What happens as we keep increasing k ?

“Optimal” value of k : $n/m \ln(2)$

- **In our case:** Optimal $k = 8 \ln(2) = 5.54 \approx 6$
 - **Error at $k = 6$:** $(1 - e^{-1/6})^2 = 0.0235$





Bloom Filter: Wrap-up

Bloom filters guarantee no false negatives, and use limited memory

- Great for pre-processing before more expensive checks

Suitable for hardware implementation

- Hash function computations can be parallelized

Options: Is it better to have **one** big **B** or **k** small **Bs**?

- It is the same: $(1 - e^{-km/n})^k$ vs. $(1 - e^{-m/(n/k)})^k$
- But keeping **1** big **B** is simpler